



KubeCon



CloudNativeCon

Europe 2019

Access Control In Kubernetes

What's Missing, And How To Fix That

@vllry

@SethMcCombs

A hand-drawn Venn diagram consisting of two overlapping circles. The larger circle, outlined in red, contains the word "Security" written in red cursive. The smaller circle, also outlined in red, is positioned inside the larger one and contains the text "Access Control" written in red cursive. The overlapping area of the two circles is shaded with light pink.

Security

Access
Control

Disclaimer



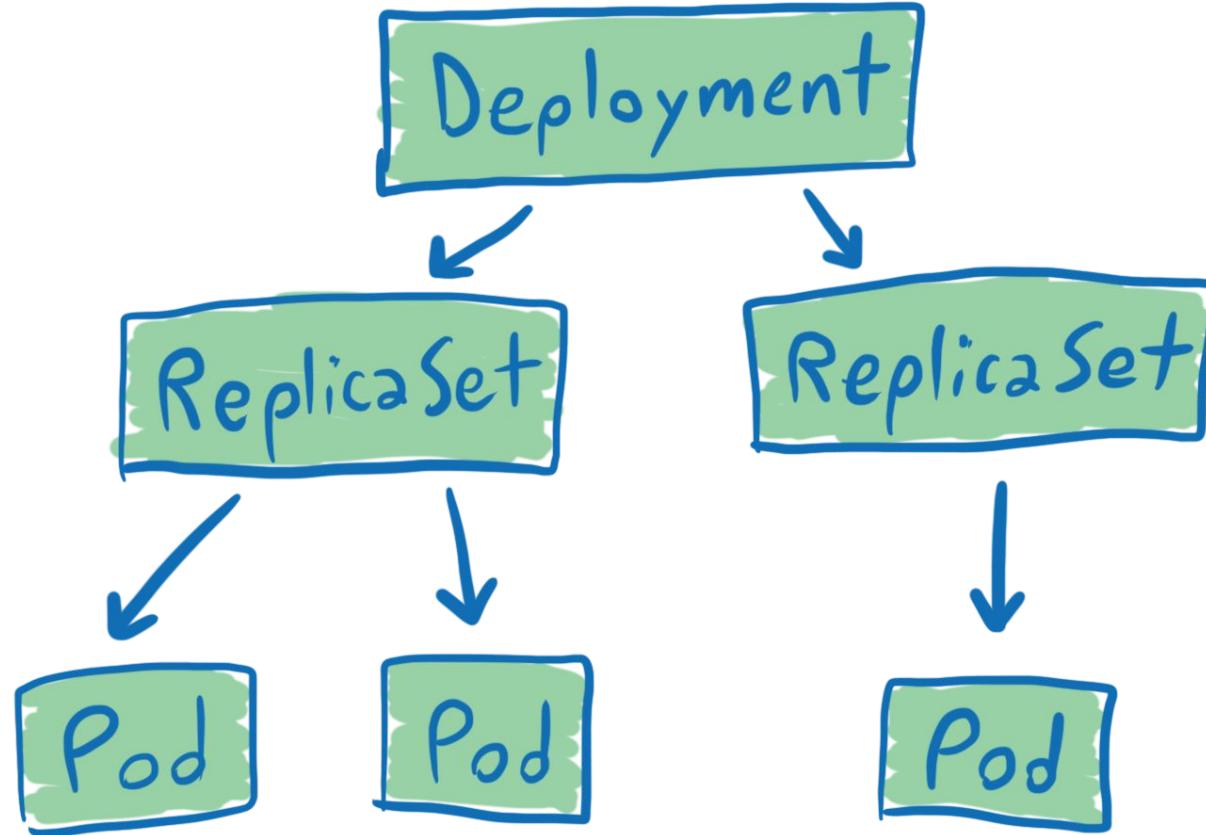
KubeCon



CloudNativeCon

Europe 2019

- This is not about Lyft systems
- This is not about Triller systems
- We might be wrong



Broad Categories



KubeCon



CloudNativeCon

Europe 2019

- Runtime characteristics
 - What can run, and in what way?
- Network access
 - What can access what?
- Meta: Kubernetes access
 - What can the app do to the Kubernetes control & data planes?

What Can We Do About All This?

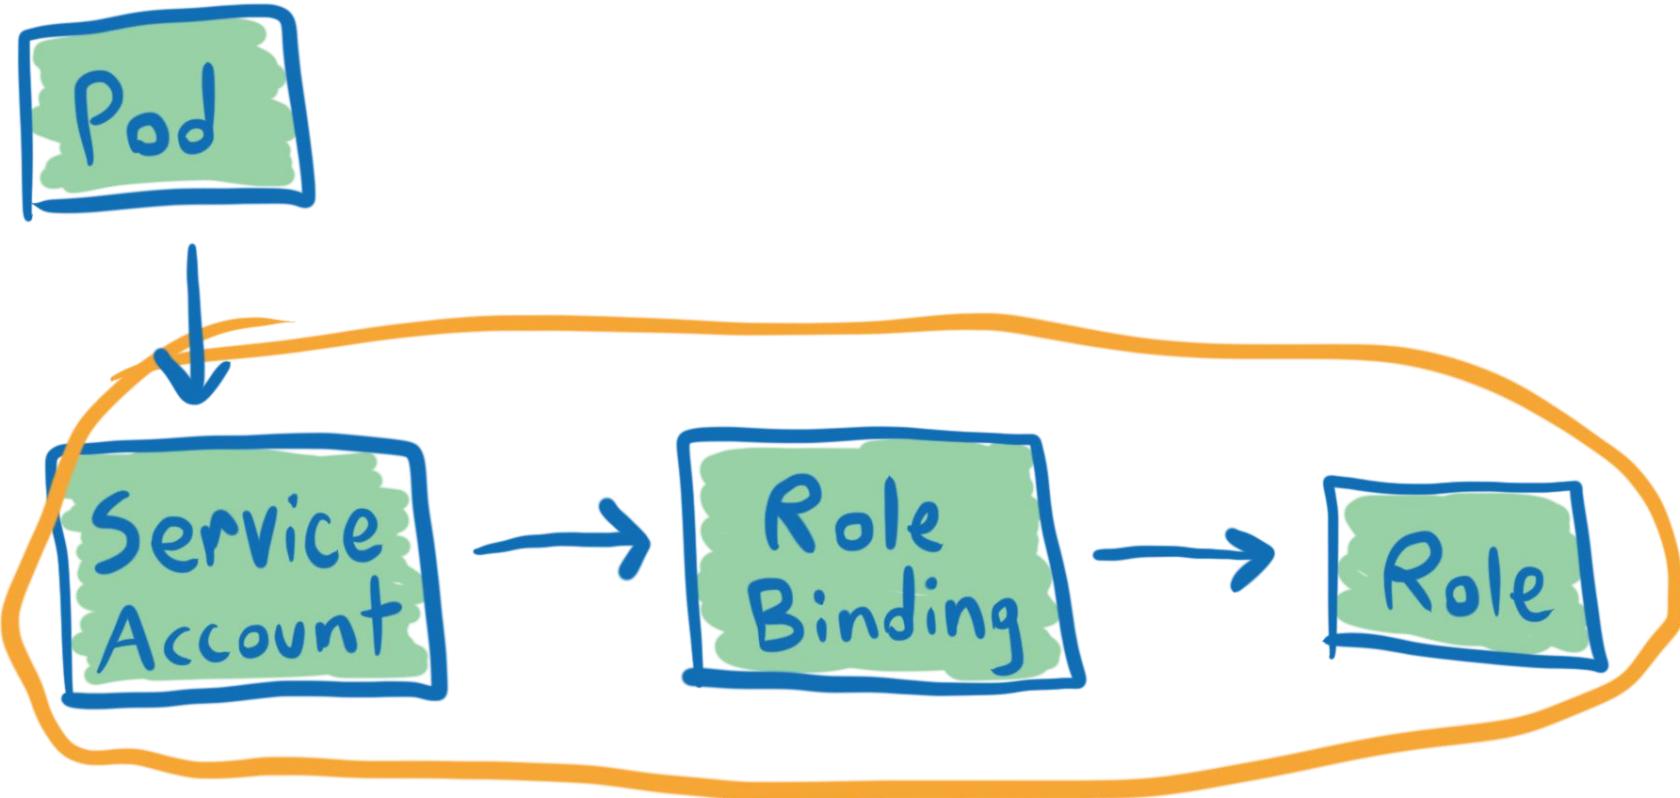
- RBAC
- Network Policies (native, Calico, Istio, etc)
- Custom API gateways
- Validating & Mutating admission webhooks (native)
- Open Policy Agent (addon)



Essentials: RBAC

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: read-pods
  namespace: default
subjects:
- kind: User
  name: vallery
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```



```
$ kubectl getrole deployments
```

DEPLOYMENTS	ROLES	PERMISSIONS
my-ingress	ingress	watch endpoints, list ingress

Other Forms of K8s API Access Control



KubeCon



CloudNativeCon

Europe 2019

- Node
 - For kubelets only. Kubelet is granted permissions based on the pods it runs.
- ABAC (old)
 - Granular user-based permissions.
- Webhook
 - Post to URL, use that response.

RBAC Limitations



KubeCon



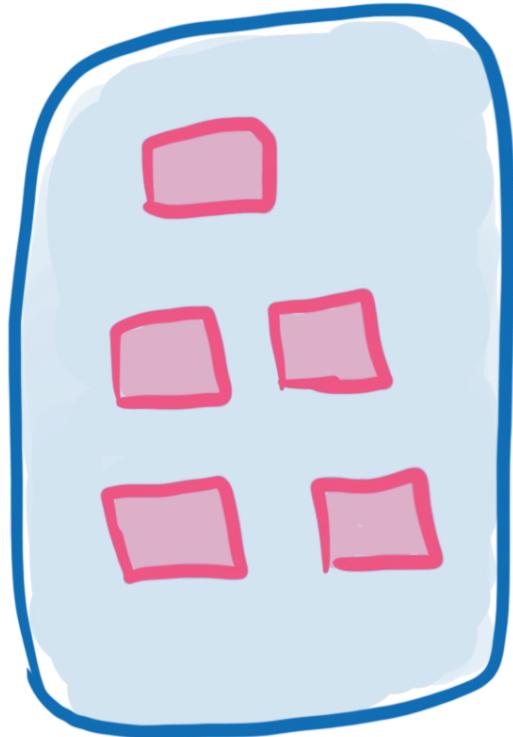
CloudNativeCon

Europe 2019

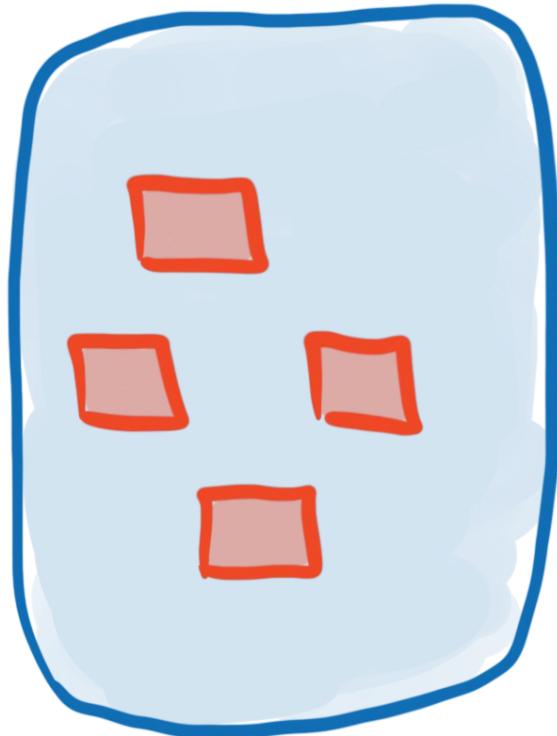
- Universal permission by resource type.
 - EG: “pods” means **all pods** in the namespace, from **any deployment**.
- No field-level access control.
 - EG: to edit the labels of an object, a user needs full write permission to the object.



Essentials: Namespaces



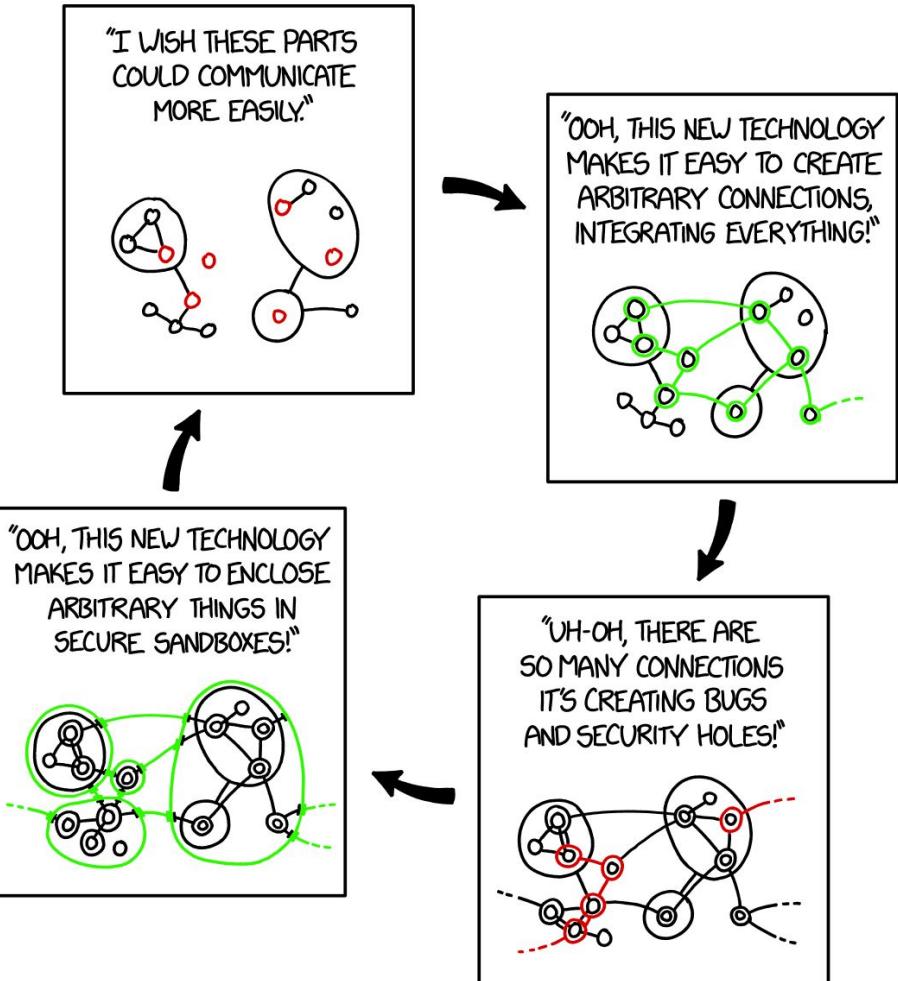
ns1



ns2



“I wish these parts could communicate more easily.”





Custom API Gateways

What Is An API Gateway?



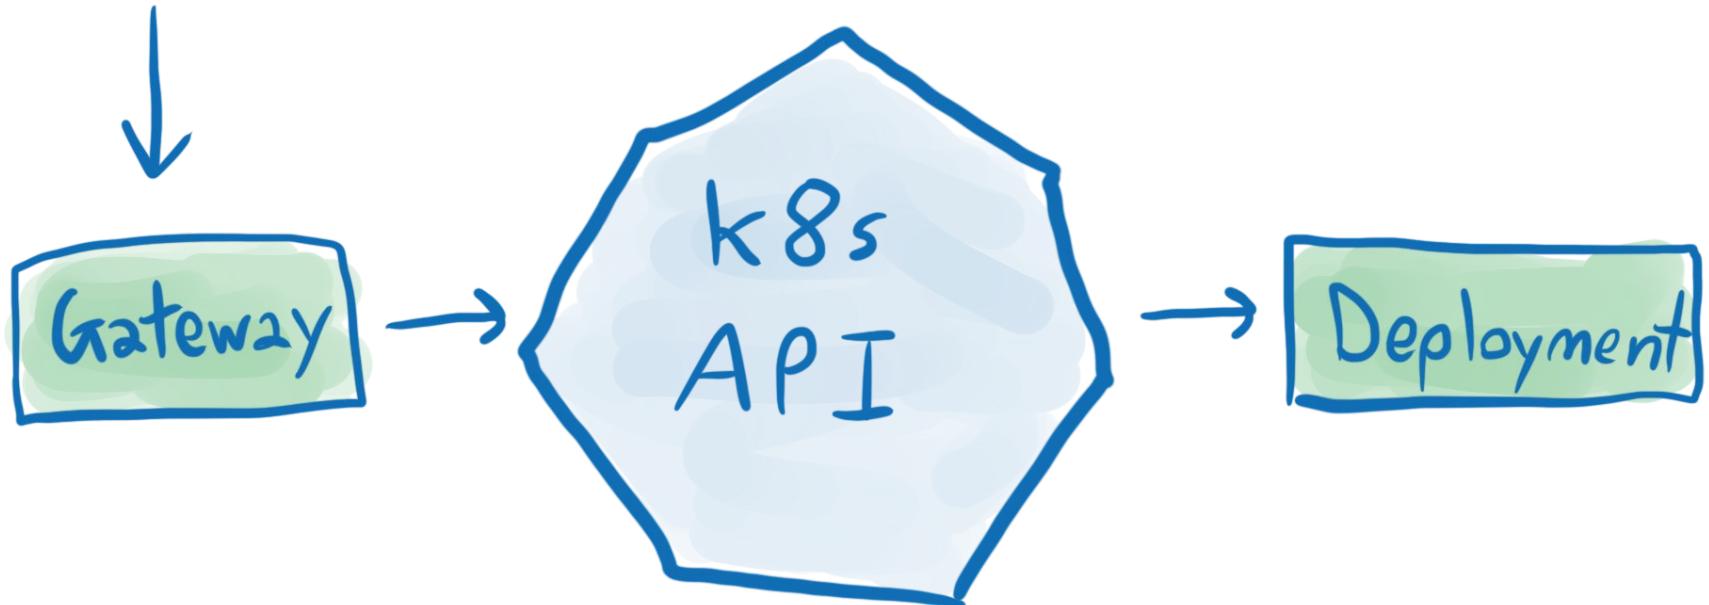
KubeCon



CloudNativeCon

Europe 2019

- A deputy/gateway is designed to perform specific actions, which require elevated permissions.
- The deputy exposes an API to trigger these actions.
- Acts as a *logical gate* to the underlying system.



API Deputy Drawbacks



KubeCon



CloudNativeCon

Europe 2019

- You still have a service with elevated permissions.
- Many actions require a very *simple* logical gate.
 - EG “only allow updates to this field”.



Admission Webhooks

Admission Webhooks



KubeCon



CloudNativeCon

Europe 2019

- Validating & Mutating WebHooks
- Custom hooks can be written to ensure resources not meeting cluster criteria are not created - Validating
- Change objects missing certain requirements - Adding labels/annotations, resource limits, etc - Mutating
- Hooks can be used to restrict the creation of objects in a Namespace based on requirements





Open Policy Agent (OPA)

Open Policy Agent

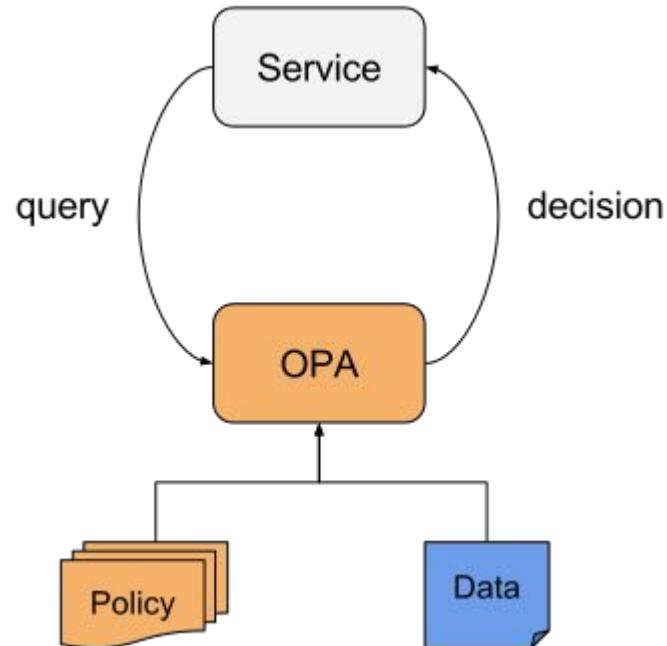


KubeCon

CloudNativeCon

Europe 2019

- OPA offloads Policy Decisions from a service (across the stack)
- Queries - can an action be taken?
 - answer is provided back to the service with an allow or deny.
- Allows context specific decisions
- Can understand the status or specs of other services



Using OPA for Resource Access Control



KubeCon



CloudNativeCon

Europe 2019

- API calls are sent to OPA with the JSON Object
- OPA compares to its data, and returns an Allow or Deny
- Universal language makes it easier to report *why* an action was denied
- Denied requests could return a patch that should be applied to the object
 - Could be used to enable a Mutating AdmissionWebhook, and apply a patch directly

Using OPA for NetworkPolicy



KubeCon



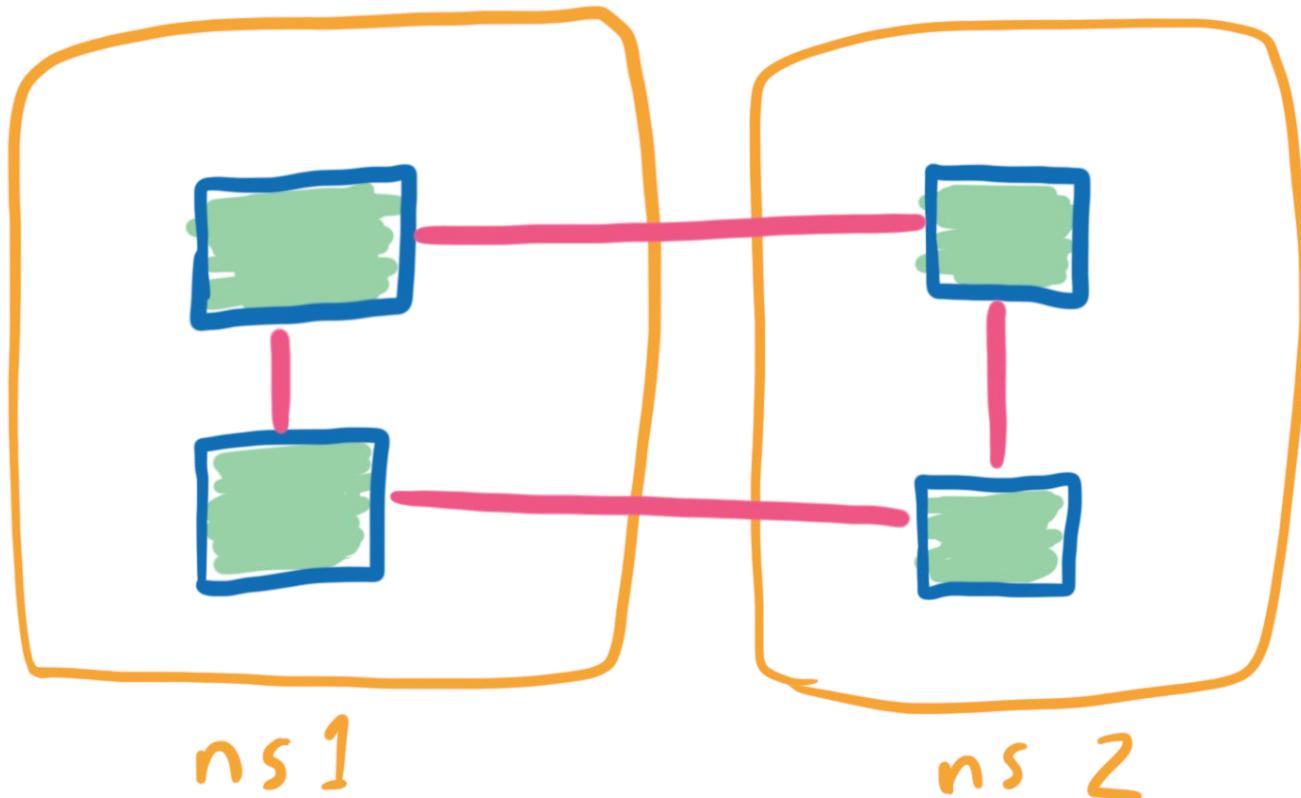
CloudNativeCon

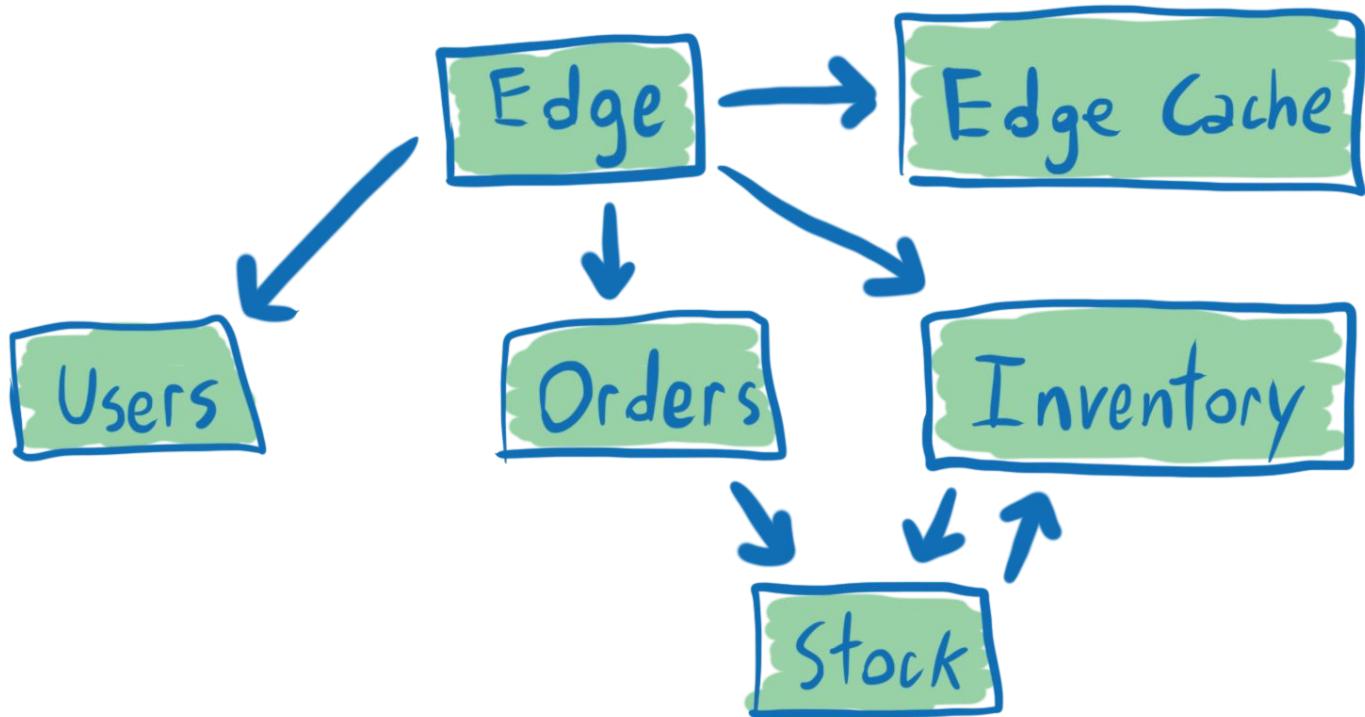
Europe 2019

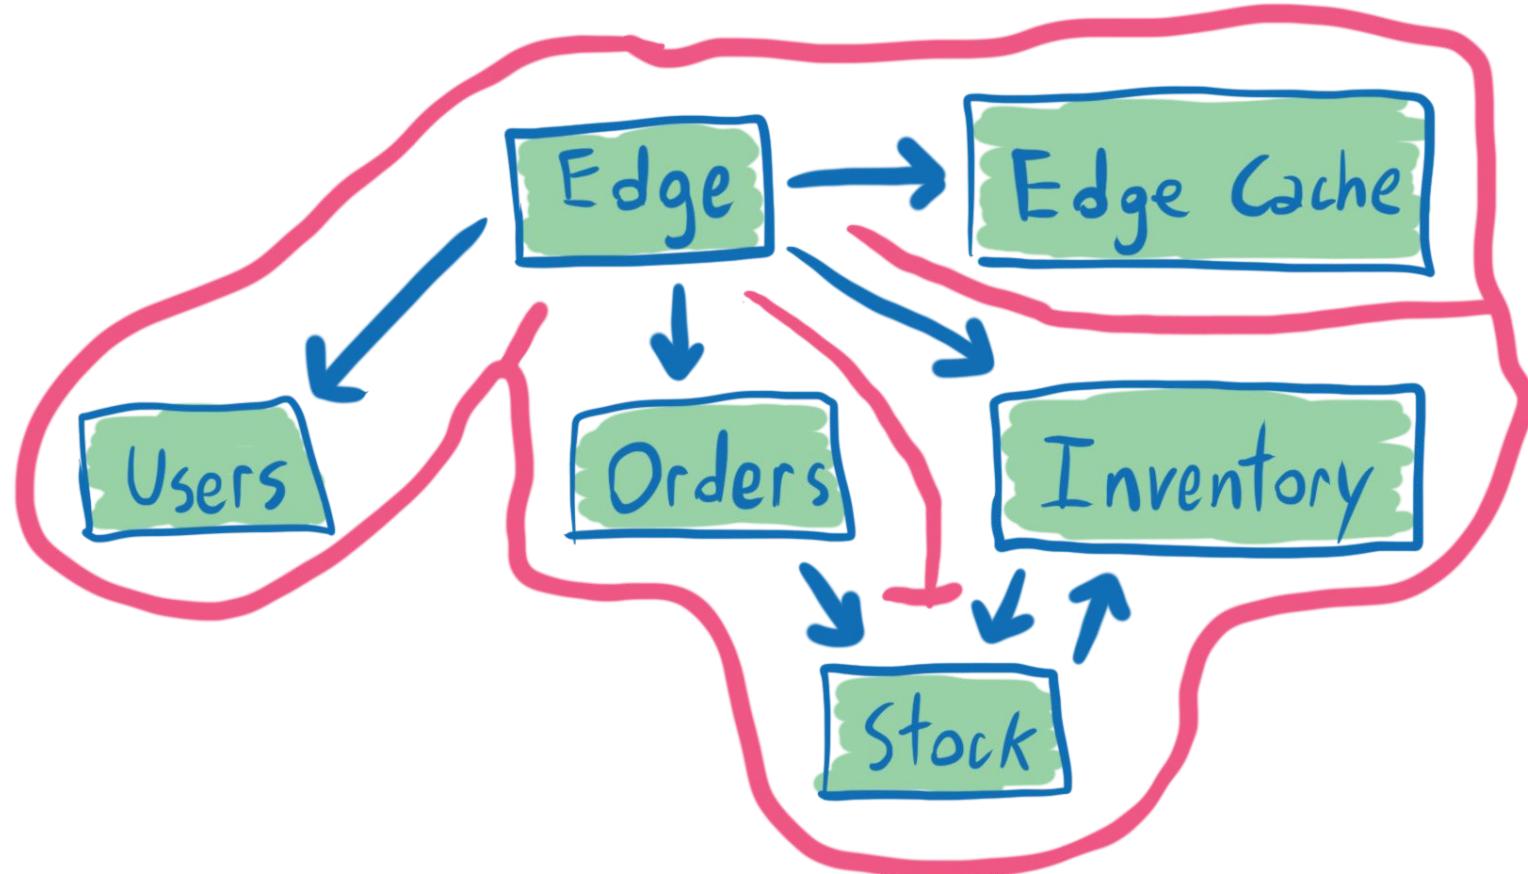
- OPA could enhance NetworkPolicy
- Layer 7 restrictions (post on /widgets but not get on /widgets)
- Labels/Annotations could be used to manage intra-Namespace communication between objects
- Powerful combo of OPA for NetworkPolicy & AdmissionControllers



Network Access Policy







Kubernetes NetworkPolicy



KubeCon



CloudNativeCon

Europe 2019

- Kubernetes has a native NetworkPolicy
 - Selectors pods with a selector label, those pods can receive no traffic unless allowed by the policy.
- Requires a 3rd-party controller - policies do nothing without one

Examples: <https://github.com/ahmetb/kubernetes-network-policy-recipes>

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: access-nginx
  namespace: demo
spec:
  podSelector:
    matchLabels:
      $KEY: $VALUE
  ingress:
    - from:
      - podSelector:
          matchLabels:
            $KEY: $VALUE
```

Implementing Network Policies



KubeCon



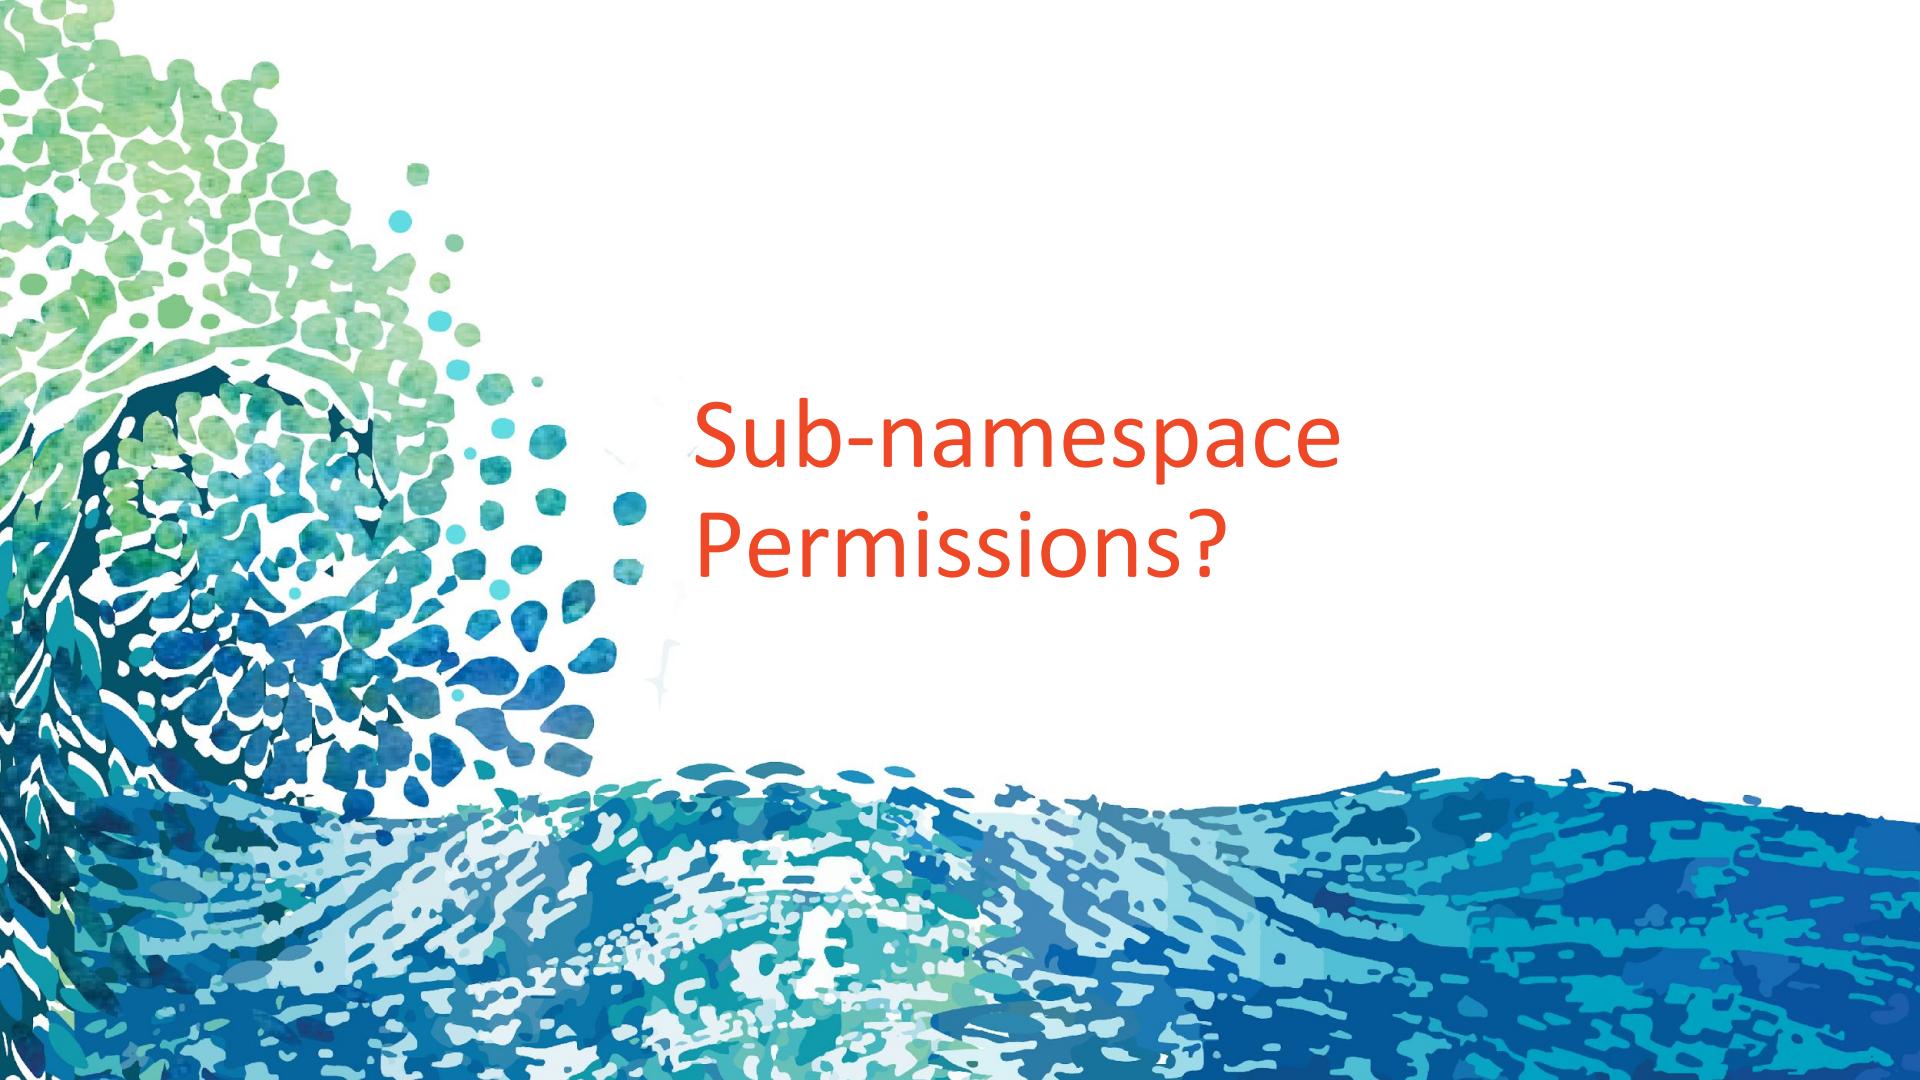
CloudNativeCon

Europe 2019

- Still an evolving addon space in Kubernetes.
 - Cilium policies
 - Istio policies
 - Calico
- Some service mesh-type tools offer similar functionality
 - Istio
 - Cilium
 - etc



Time For The Opinions!



Sub-namespace Permissions?

```
SpaceshipGrey:~ vallery$ ls -l
total 32
drwx-----@    5 vallery  staff    ... Applications
drwx-----+  17 vallery  staff    ... Desktop
```

apiVersion: ...

kind: ...

metadata:

...

spec:

...

status:

...

apiVersion: ...

kind: ...

metadata:

owningRoles:

- jenkins
- ops

...

spec:

...

status:

...





Smaller Namespaces?

What Are Our Limiting Factors?



KubeCon



CloudNativeCon

Europe 2019

- Objects that rely on one another need to be in the same namespace.
 - EG Ingress / Service / Deployment / Pods
- (Namespace level) service accounts can't be used in multiple namespaces.
 - Users / bots will need more accounts.

Cookbook For Success



KubeCon



CloudNativeCon

Europe 2019

- Keep ~1 distinct (lowercase s) service per namespace.
- Restrict pod traffic between namespaces with a NetworkPolicy, to only allow required cross-namespace traffic.
- Use multiple controllers/SAs to operate on objects spread between namespaces.
 - Use a cluster-level controller/SA if the scope/access is appropriate.



Too Long, Didn't Listen

0. Kubernetes' design priority is extensibility, not out-of-the-box usability. This includes security considerations.

1. Put individual sets of resources, like a service/project in their own namespaces.

2. Use RBAC to restrict access to **reasonable** as-needed access. Prefer too-broad over painfully granular.

3. Use policy tools like webhooks or OPA to provide narrow restrictions to resources, beyond RBAC granularity.

Vallery Lancey

- Infrastructure at Lyft
- @vllry on Twitter and Github

Kubernetes contributor (mostly SIG-network). Dabbles in many aspects of distributed systems and reliability engineering.

Seth McCombs

- SRE at Triller
- Tweets at @SethMcCombs

More ops than dev, fiddles with containers and Kubernetes. Avid guitarist and collector of fountain pens.